

```
Model: "sequential_27"
```

Layer (type)	Output Shape	Param #
conv2d_105 (Conv2D)	(None, 26, 26, 64)	640
max_pooling2d_62 (MaxPooling2D)	(None, 13, 13, 64)	0
conv2d_106 (Conv2D)	(None, 11, 11, 128)	73,856
max_pooling2d_63 (MaxPooling2D)	(None, 5, 5, 128)	0
conv2d_107 (Conv2D)	(None, 3, 3, 256)	295,168
conv2d_108 (Conv2D)	(None, 1, 1, 512)	1,180,160
flatten_19 (Flatten)	(None, 512)	0
dense_38 (Dense)	(None, 128)	65,664
dense_39 (Dense)	(None, 10)	1,290

```
Total params: 1,616,770 (6.17 MB)
Trainable params: 1,616,770 (6.17 MB)
Non-trainable params: 0 (0.00 B)
Epoch 1/5
844/844 — 11s 9ms/step - accuracy: 0.8947 - loss: 0.3310 - val_accuracy: 0.9870 - val_loss: 0.0463
Epoch 2/5
844/844 — 5s 6ms/step - accuracy: 0.9870 - loss: 0.0422 - val_accuracy: 0.9890 - val_loss: 0.0387
Epoch 3/5
844/844 — 5s 6ms/step - accuracy: 0.9909 - loss: 0.0298 - val_accuracy: 0.9878 - val_loss: 0.0401
Epoch 4/5
844/844 — 5s 6ms/step - accuracy: 0.9928 - loss: 0.0231 - val_accuracy: 0.9915 - val_loss: 0.0335
Epoch 5/5
844/844 — 5s 6ms/step - accuracy: 0.9946 - loss: 0.0165 - val_accuracy: 0.9913 - val_loss: 0.0371
y_test shape (before one-hot encoding): (10000,)
y_test shape (after one-hot encoding): (10000, 10)
313/313 — 1s 3ms/step - accuracy: 0.9884 - loss: 0.0461
Test accuracy: 0.9916, Test loss: 0.0332
```

Source Code:

```
from tensorflow.keras.datasets import mnist
import cv2
from google.colab.patches import cv2_imshow
from matplotlib import pyplot as plt
import numpy as np
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, Dense, Flatten, MaxPool2D

(X_train, y_train), (X_test, y_test) = mnist.load_data()

X_train = X_train.reshape(60000, 28, 28, 1)
X_test = X_test.reshape(10000, 28, 28, 1)

X_train = X_train.astype("float32") / 255
X_test = X_test.astype("float32") / 255

print(y_train)
print(y_train.shape)
print(y_train[0])
y_train = to_categorical(y_train)
print(f'categorical data:\n{y_train}')
print(y_train.shape)
print(y_train[0])

cnn = Sequential()
cnn.add(Conv2D(filters=64, kernel_size=(3, 3), activation="relu", input_shape=(28,28, 1)))
cnn.add(MaxPool2D(pool_size=(2, 2)))
cnn.add(Conv2D(filters=128, kernel_size=(3, 3), activation="relu"))
cnn.add(MaxPool2D(pool_size=(2, 2)))

#add CNN layers
cnn.add(Conv2D(filters=256, kernel_size=(3, 3), activation="relu"))
cnn.add(Conv2D(filters=512, kernel_size=(3, 3), activation="relu"))

cnn.add(Flatten())
cnn.add(Dense(units=128, activation="relu"))
cnn.add(Dense(units=10, activation="softmax"))

cnn.summary()
```

```

cnn.compile(optimizer="adam", loss="categorical_crossentropy", metrics=["accuracy"])

cnn.fit(X_train, y_train, epochs=5, batch_size=64, validation_split=0.1)

~
# Evaluate accuracy on test dataset...
# Convert y_test to one-hot encoded format
print(f'y_test shape (before one-hot encoding): {y_test.shape}')
y_test = to_categorical(y_test) # convert to one-hot encoded labels (shape (None,10))
print(f'y_test shape (after one-hot encoding): {y_test.shape}')
# ...Evaluate accuracy on test dataset
test_loss, test_accuracy = cnn.evaluate(X_test, y_test)
print(f"Test accuracy: {test_accuracy:.4f}, Test loss: {test_loss:.4f}")

# After defining the model (cnn), print weights for Conv2D layers
for i, layer in enumerate(cnn.layers):
    if isinstance(layer, Conv2D): # Check if the layer is Conv2D
        weights, biases = layer.get_weights()
        print(f"Conv2D Layer {i} Weights (Filters):")
        print(weights.shape) # Shape of the weight tensor
        print(weights)
        print(f"Conv2D Layer {i} Biases:")
        print(biases.shape)
        print(biases)

```


Exercise 2

In continuation of Ex. 1, **predict** the classification outcome of two handwritten digits. Print the images as well as the predictions.


Output:






1/1  0s 31ms/step
The prediction is: 3




1/1  0s 28ms/step
The prediction is: 2



1/1  0s 30ms/step
The prediction is: 9



1/1  0s 30ms/step
The prediction is: 4

Source Code:

```

from tensorflow.keras.datasets import mnist
import cv2
from google.colab.patches import cv2_imshow
from matplotlib import pyplot as plt
import numpy as np
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, Dense, Flatten, MaxPool2D

(X_train, y_train), (X_test, y_test) = mnist.load_data()

X_train = X_train.reshape(60000, 28, 28, 1)
X_test = X_test.reshape(10000, 28, 28, 1)

X_train = X_train.astype("float32") / 255
X_test = X_test.astype("float32") / 255

print(y_train)
print(y_train.shape)
print(y_train[0])
y_train = to_categorical(y_train)
print(f'categorical data:\n{y_train}')
print(y_train.shape)
print(y_train[0])

cnn = Sequential()

cnn.add(Conv2D(filters=64, kernel_size=(3, 3), activation="relu", input_shape=(28,28, 1)))
cnn.add(MaxPool2D(pool_size=(2, 2)))
cnn.add(Conv2D(filters=128, kernel_size=(3, 3), activation="relu"))
cnn.add(MaxPool2D(pool_size=(2, 2)))

#add CNN layers
cnn.add(Conv2D(filters=256, kernel_size=(3, 3), activation="relu"))
cnn.add(Conv2D(filters=512, kernel_size=(3, 3), activation="relu"))

cnn.add(Flatten())
cnn.add(Dense(units=128, activation="relu"))
cnn.add(Dense(units=10, activation="softmax"))

cnn.summary()

cnn.compile(optimizer="adam", loss="categorical_crossentropy",
metrics=["accuracy"])

history = cnn.fit(X_train, y_train, epochs=5, batch_size=64, validation_split=0.1)

# Evaluate accuracy on test dataset...
# Convert y_test to one-hot encoded format
print(f'y_test shape (before one-hot encoding): {y_test.shape}')
y_test = to_categorical(y_test) # convert to one-hot encoded labels (shape (None,10)
print(f'y_test shape (after one-hot encoding): {y_test.shape}')

```

```

# ...Evaluate accuracy on test dataset
test_loss, test_accuracy = cnn.evaluate(X_test, y_test)
print(f"Test accuracy: {test_accuracy:.4f}, Test loss: {test_loss:.4f}")

# After defining the model (cnn), print weights for Conv2D layers
for i, layer in enumerate(cnn.layers):
    if isinstance(layer, Conv2D): # Check if the layer is Conv2D
        weights, biases = layer.get_weights()
        print(f"Conv2D Layer {i} Weights (Filters):")
        print(weights.shape) # Shape of the weight tensor
        print(weights)
        print(f"Conv2D Layer {i} Biases:")
        print(biases.shape)
        print(biases)

def myPredict(image):
    image = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
    image = cv2.resize(image, (28, 28))
    image = image.reshape((1, 28, 28, 1))

    # MNIST images have dark background and light foreground,
    # we need to correct our images to match that
    image = cv2.bitwise_not(image)
    prediction = cnn.predict(image)[0]
    return np.argmax(prediction)

#5
image = cv2.imread('sample_data/five.jpg')
cv2.imshow(image)
print(f'The prediction is: {myPredict(image)}')

#3
image = cv2.imread('sample_data/three.png')
cv2.imshow(image)
print(f'The prediction is: {myPredict(image)}')

#2
image = cv2.imread('sample_data/two.png')
cv2.imshow(image)
print(f'The prediction is: {myPredict(image)}')

#new predictions
#9
image = cv2.imread('sample_data/nine.png')
cv2.imshow(image)
print(f'The prediction is: {myPredict(image)}')

#4
image = cv2.imread('sample_data/four.png')
cv2.imshow(image)
print(f'The prediction is: {myPredict(image)}')

```

Exercise 3

Given the following datasets: train.zip and validation.zip, and the accompanying source code: Lec37.1_CNN_cats+dogs.py. This a CNN that classifies images between cats and dogs. Modify the number of **EPOCHS** to a larger value. Report whether the accuracy improves. In addition, **predict** the classification outcome of two images, one **cat** and one **dog**. Print the images as well as the predictions. The two images should not be from the datasets.

Note: The code should run without any issues on Visual Studio.

Output:



```
← → + Q ≡ 📁
history.hCat Image Prediction: Cat
= histoDog Image Prediction: Dog
ppgs =
```

Source Code:


```

1  import tensorflow as tf
2  from tensorflow.keras.preprocessing.image import ImageDataGenerator, load_img, img_to_array
3  from tensorflow.keras.models import Sequential
4  from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
5  import os
6  import matplotlib.pyplot as plt
7  import numpy as np
8  import cv2
9
10 # Define paths to the dataset
11 train_dir = 'train'
12 validation_dir = 'validation'
13
14 # Image parameters
15 IMG_HEIGHT = 150
16 IMG_WIDTH = 150
17 BATCH_SIZE = 32
18
19 # Data augmentation and preprocessing for training
20 train_datagen = ImageDataGenerator(
21     rescale=1./255,
22     rotation_range=40,
23     width_shift_range=0.2,
24     height_shift_range=0.2,
25     shear_range=0.2,
26     zoom_range=0.2,
27     horizontal_flip=True,
28     fill_mode='nearest'
29 )
30
31 # Only rescaling for validation
32 validation_datagen = ImageDataGenerator(rescale=1./255)
33
34 # Load and preprocess training data
35 train_generator = train_datagen.flow_from_directory(
36     train_dir,
37     target_size=(IMG_HEIGHT, IMG_WIDTH),
38     batch_size=BATCH_SIZE,
39     class_mode='binary' # Binary classification: cats (0) vs dogs (1)
40 )
41
42 # Load and preprocess validation data
43 validation_generator = validation_datagen.flow_from_directory(
44     validation_dir,
45     target_size=(IMG_HEIGHT, IMG_WIDTH),
46     batch_size=BATCH_SIZE,
47     class_mode='binary'
48 )
49

```

```

50     # Build the CNN model
51     model = Sequential([
52         Conv2D(32, (3, 3), activation='relu', input_shape=(IMG_HEIGHT, IMG_WIDTH, 3)),
53         MaxPooling2D(2, 2),
54         Conv2D(64, (3, 3), activation='relu'),
55         MaxPooling2D(2, 2),
56         Conv2D(128, (3, 3), activation='relu'),
57         MaxPooling2D(2, 2),
58         Flatten(),
59         Dense(512, activation='relu'),
60         Dropout(0.5),
61         Dense(1, activation='sigmoid') # Binary output
62     ])
63
64     # Compile the model
65     model.compile(optimizer='adam',
66                  loss='binary_crossentropy',
67                  metrics=['accuracy'])
68
69     # Model summary
70     model.summary()
71
72     # Train the model
73     EPOCHS = 5
74     history = model.fit(
75         train_generator,
76         steps_per_epoch=train_generator.samples // BATCH_SIZE,
77         epochs=EPOCHS,
78         validation_data=validation_generator,
79         validation_steps=validation_generator.samples // BATCH_SIZE
80     )
81
82     # Evaluate the model
83     loss, accuracy = model.evaluate(validation_generator)
84     print(f"Validation Loss: {loss:.4f}")
85     print(f"Validation Accuracy: {accuracy:.4f}")
86
87     # Save the model
88     model.save('cats_vs_dogs_model.h5')
89
90     acc = history.history['accuracy']
91     val_acc = history.history['val_accuracy']
92     loss = history.history['loss']
93     val_loss = history.history['val_loss']
94     epochs_range = range(EPOCHS)
95

```

```

96     plt.figure(figsize=(12, 4))
97     plt.subplot(1, 2, 1)
98     plt.plot(epochs_range, acc, label='Training Accuracy')
99     plt.plot(epochs_range, val_acc, label='Validation Accuracy')
100    plt.title('Training and Validation Accuracy')
101    plt.legend()
102
103    plt.subplot(1, 2, 2)
104    plt.plot(epochs_range, loss, label='Training Loss')
105    plt.plot(epochs_range, val_loss, label='Validation Loss')
106    plt.title('Training and Validation Loss')
107    plt.legend()
108
109    plt.savefig('training_history.png')
110    plt.close()
111
112    def predict_image(image_path):
113        img = load_img(image_path, target_size=(IMG_HEIGHT, IMG_WIDTH))
114        img_array = img_to_array(img)
115        img_array = np.expand_dims(img_array, axis=0)
116        img_array /= 255.0
117        prediction = model.predict(img_array)
118        return prediction
119
120    cat_prediction = predict_image('cat_pic.png')
121    dog_prediction = predict_image('dog_pic.png')
122    cat_img = cv2.imread('cat_pic.png')
123    dog_img = cv2.imread('dog_pic.png')
124    print(f"Cat Image Prediction: {'Dog' if cat_prediction > 0.5 else 'Cat'}")
125    print(f"Dog Image Prediction: {'Dog' if dog_prediction > 0.5 else 'Cat'}")
126    plt.figure(figsize=(10, 5))
127    plt.subplot(1, 2, 1)
128    plt.imshow(cv2.cvtColor(cat_img, cv2.COLOR_BGR2RGB))
129    plt.axis('off')
130
131    plt.subplot(1, 2, 2)
132    plt.imshow(cv2.cvtColor(dog_img, cv2.COLOR_BGR2RGB))
133    plt.axis('off')
134    plt.show()

```

Note: Submit through Canvas